

Markov chain Monte Carlo methods for Uncertainty Quantification

General concept

Lecture 21

March 31, 2016; April 5, 2016; April 7, 2016

Motivation: Bayesian inference & prediction

Data: $D = (D_1, \dots, D_N)$

Parameters: $X = (X_1, \dots, X_d)$

Likelihood: $\mathcal{L}(D|X)$

Prior model: $\pi(X)$

Posterior model: The basis for inference about X
... via Bayes theorem

$$\pi(X|D) = \frac{\mathcal{L}(D|X)\pi(X)}{\int \mathcal{L}(D|Y)\pi(Y)dY}$$

Predictive model: Predictive distribution of a future obs. D_{future} .

$$\mathcal{L}(D_{\text{future}}|D) = \int \mathcal{L}(D_{\text{future}}|X)\pi(X|D)dX$$

... expected likelihood where uncertainty of X is
constrained w.r.t $\pi(X|D)$

The problem: Challenges in Bayesian Inference

For some function $g(\cdot)$,

- ▶ the derivation of any posterior quantity requires the computation of integrals of the form:

$$E_{\pi(X|D)}(g(X)) = \int g(X)\pi(X|D)dX$$

- ▶ the posterior distribution density $\pi(X|D)$ or $\pi(g(X)|D)$ is intractable because of

$$\int \mathcal{L}(D|X)\pi(X)dX = ??$$

CODE

CODE

Application: Bayesian hierarchical model

Likelihood

$$D_i \sim \text{Bernoulli}(p(t_i|\alpha, \beta)), i = 1, \dots, 23$$

$$p(t_i|\alpha, \beta) = \frac{\exp(\alpha + \beta t_i)}{1 + \exp(\alpha + \beta t_i)}$$

Prior

$$\alpha \sim \text{N}(\mu = 0, \sigma^2 = 10^2)$$

$$\beta \sim \text{N}(\mu = 0, \sigma^2 = 10^2)$$

Posterior:

$$\begin{aligned} \pi(X = (\alpha, \beta) | D) &= \prod_{i=1}^{23} \left(\frac{\exp(\alpha + \beta t_i)}{1 + \exp(\alpha + \beta t_i)} \right)^{D_i} \left(\frac{1}{1 + \exp(\alpha + \beta t_i)} \right)^{1-D_i} \\ &\quad \times \exp\left(-\frac{1}{2}\alpha^2/10^2\right) \times \exp\left(-\frac{1}{2}\beta^2/10^2\right) \times \frac{1}{\text{CONST.}} \end{aligned}$$

The cure: Monte Carlo methods

Due to the 'essential correspondence' between
density $\pi(X|D)$ & samples $\{X^{(n)} \sim \pi(X|D)\}$:

- ▶ Posterior density could be re-created via
 - ▶ histograms estimators,
 - ▶ kernel density estimators,
 - ▶ Normal mixture models, etc...
- ▶ Expectations could be approx. via

$$E_{\pi(X|D)}(g(X)) \approx \frac{1}{N} \sum_{n=1}^N g(x^{(n)})$$

Monte Carlo methods (main idea)

- ▶ Generate an i.i.d. sample $X^{(n)} \sim \pi(dX)$, for $n = 1, \dots, N$
 - ▶ Inverse probability integral transform
 - ▶ Rejection sampling
 - ▶ Importance sampling

- ▶ Approx. integral $E_{\pi}(g(X)) = \int g(X)\pi(X)dX$

$$\text{with } \bar{g}^{(N)} \approx \frac{1}{N} \sum_{n=1}^N g(X^{(n)})$$

$$\text{and standard error s.e.}(\bar{g}^{(N)}) = \sqrt{\frac{1}{N} \text{Var}_{\pi}(g(X))}$$

... according to the \sqrt{N} -CLT

Ripley (2001). Stochastic simulation.

Markov chain Monte Carlo methods (main idea)

- ▶ Generate a Markov chain $X^{(n)} \sim P(d \cdot | X^{(n-1)})$, for $n = 1, \dots, N$

- ▶ Approx. integral $E_{\pi}(g(X)) = \int g(X)\pi(X)dX$

$$\text{with } \bar{g}^{(N)} \approx \frac{1}{N} \sum_{n=1}^N g(X^{(n)})$$

$$\text{and standard error s.e.}(\bar{g}^{(N)}) = \sqrt{\frac{1}{N} \tau_g \text{Var}_{\pi}(g(X))}$$

where $\tau_g \in (0, \infty)$ is the integrated autocorrelation time with $\tau_g = 1 + 2 \sum_{k=1}^{\infty} \text{Cor}(x_n, x_{n+k})$

... according to the (Markov chain) \sqrt{N} -CLT

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

MCMC theory

Main conditions for $P(d \cdot | \cdot)$

1. Stationarity w.r.t. $\pi(d \cdot)$
 - Reversibility w.r.t. $\pi(d \cdot)$
2. ϕ -irreducibility
 - Harris recurrent
3. Aperiodicity

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

Stationarity

Definition: The Markov chain $\{x^{(n)}; n = 1, \dots, N\}$ simulated by the transition probability $P(d \cdot | \cdot)$ has stationary (or invariant) distribution $\pi(\cdot)$ iff

$$\int_{x \in \mathcal{X}} \pi(dx) P(dy|x) = \pi(dy)$$

where $x, y \in \mathcal{X}$.

Explanation: If $x_n \sim \pi(d \cdot)$ & $x_{n+1} \sim P(d \cdot | x_n)$, then $x_{n+1} \sim \pi(d \cdot)$

Hopefully, if we run the Markov chain (started from anywhere) for a long time, then for a long N the distribution of X_N will be approx. stationary:

$$X_N \stackrel{\text{approx.}}{\sim} \pi(d \cdot).$$

Reversibility

Definition: The Markov chain $\{x^{(n)}; n = 1, \dots\}$ simulated by the transition probability $P(d \cdot | \cdot)$ is reversible w.r.t distribution $\pi(\cdot)$ iff

$$\pi(dx)P(dy|x) = \pi(dy)P(dx|y)$$

where $x, y \in \mathcal{X}$.

Explanation: It expresses an equilibrium in the flow of the Markov chain: The probability of being in x and moving to y is the same as the probability of being in y and moving to x .

Property: Reversibility implies stationarity

Rational: It is more conservative assumption, but it is easier to be checked, since no integral is involved

Reversibility implies stationarity

Proposition: If Markov chain $\{x^{(n)}\}$ with transition probability $P(d \cdot | \cdot)$ is reversible w.r.t distribution $\pi(d \cdot)$, then $\pi(d \cdot)$ is the stationarity distr.

Prof: We compute that

$$\begin{aligned}\int_{x \in \mathcal{X}} \pi(dx) P(dy|x) &= \int_{x \in \mathcal{X}} \pi(dy) P(dx|y) \\ &= \pi(dy) \int_{x \in \mathcal{X}} P(dx|y) \\ &= \pi(dy)\end{aligned}$$

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

ϕ -irreducibility

Definition: The Markov chain $\{x^{(n)}\}$ with transition probability $P(d \cdot | \cdot)$ is ϕ -irreducible if for all $A \subseteq \mathcal{X}$ with $\phi(A) > 0$, there exists a positive integer n s.t. $P^n(A|x) > 0$, for all $x \in \mathcal{X}$, where $P^n(A|x)$ is the n -step transition probability of the Markov chain.

Explain: The Markov chain has positive probability of eventually reaching any state from any other state, in a finite number of iterations.

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

Harris recurrent

Definition: The Markov chain $\{x^{(n)}\}$ with transition probability $P(d \cdot | \cdot)$ is Harris recurrent if for all $A \subseteq \mathcal{X}$ with $\pi(A) > 0$ and for all $x \in \mathcal{X}$, there exists a positive integer n s.t. $P^n(A|x) = 1$, for all $x \in \mathcal{X}$, where $P^n(A|x)$ is the n -step transition probability of the Markov chain.

Explain: For all $A \subseteq \mathcal{X}$ with $\pi(A) > 0$ and for all $x \in \mathcal{X}$, the probability that the Markov chain will eventually reach B from x is 1.

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

Aperiodicity

The Markov chain $\{x^{(n)}\}$ with transition probability $P(d \cdot | \cdot)$ is aperiodic if there does not exist partition $\{\mathcal{X}_i; i = 1, \dots, \kappa\}$ where $\pi(\mathcal{X}_i) > 0$ s.t.

- ▶ $P(\mathcal{X}_{i+1}|x) = 1$ for $x \in \mathcal{X}_i$ and
- ▶ $P(\mathcal{X}_1|x) = 1$ for $x \in \mathcal{X}_\kappa$

Christian P. Robert and George Casella (2004). Monte Carlo Statistical Methods.

Ergodicity

Theorem: If Markov chain $\{x^{(t)}\}$ with transition probability $P(d \cdot | \cdot)$ is Harris recurrent, aperiodic, and has a stationary distribution $\pi(d \cdot)$ then for every initial distribution $\tilde{\pi}$

$$\lim_{n \rightarrow \infty} \left\| \int P^n(\cdot | x) \tilde{\pi}(dx) - \pi(\cdot) \right\| = 0$$

Explain: Standard Markov chain theory tells that for any initial seed $x^{(0)}$, the realisation of the chain $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots\}$, provides via the ergodic theorem, a realisation of the stationary distribution since

$$x^{(n)} \rightarrow \pi(\cdot), \text{ as } n \rightarrow \infty$$

Markov chain \sqrt{N} -CLT

Theorem: If Markov chain $\{x^{(n)}\}$ with transition probability $P(d \cdot | \cdot)$ is irreducible, aperiodic, and reversible with stationary distribution $\pi(d \cdot)$ then the CLT applies:
For some function $g(\cdot)$, and $\bar{g}^{(N)} = \frac{1}{N} \sum_{n=1}^N g(x^{(n)})$

$$N^{-1/2}(\bar{g}^{(N)} - E_{\pi}(g(x))) \implies N(0, \tau_g \text{Var}_{\pi}(g(x)))$$

where $\tau_g = 1 + 2 \sum_{k=1}^{\infty} \text{Cor}(x_n, x_{n+k})$, if $\tau_g < \infty$.

Explain: Standard Markov chain theory tells that for any initial seed $x^{(0)}$, the realisation of the chain $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots\}$, provides, an approx. of the required expectations as

$$\bar{g}^{(N)} \rightarrow E_{\pi}(g(x)), \text{ as } N \rightarrow \infty$$

where $\bar{g}^{(N)} = \frac{1}{N} \sum_{n=1}^N g(x^{(n)})$.

CODE

CODE

Metropolis-Hastings: the algorithm

Simulate from a Metropolis-Hastings transition probability $P(d \cdot | \cdot)$, with target distribution $\pi(d \cdot)$, and proposal distribution $q(d \cdot | \cdot)$.

i.e. $X^{(n+1)} \sim P(d \cdot | X^{(n)})$

Given that the current state of the Markov chain is at $X^{(n)} = x$:

1. Generate a 'proposed value' x' from $q(d \cdot | x)$
2. Calculate

$$a(x, x') = \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x|x')}{q(x'|x)}\right)$$

3. With probability $a(x, x')$ accept the proposed value and set $X^{(n+1)} = x'$; otherwise reject and set $X^{(n+1)} = x$.

Generate $u \sim U(0, 1)$, and set $X^{(n+1)} = \begin{cases} x' & , \text{ if } a(x, x') \geq u \\ x & , \text{ if } a(x, x') < u \end{cases}$

Metropolis-Hastings: the transition probability

The Metropolis-Hastings transition probability is:

$$P(y|x) = q(y|x)a(x,y) + (1 - r(x))\delta_x(y)$$

where $r(x) = \int q(y|x)a(x,y)dy$,

and $\delta_x(y)$ is the Dirac mass in x .

Metropolis-Hastings: Reversibility

The Metropolis-Hastings (as described above) produces a Markov chain $\{x^{(t)}\}$ which is reversible w.r.t $\pi(d\cdot)$.

Prof: We need to show that

$$\pi(dx)P(dy|x) = \pi(dy)P(dx|y)$$

If $x \neq y$,

$$\begin{aligned}\pi(dx)P(dy|x) &= [\pi(x)dx][q(y|x)a(x,y)dy] \\ &= \pi(x)q(y|x) \min(1, \frac{\pi(y)}{\pi(x)} \frac{q(x|y)}{q(y|x)}) dx dy \\ &= \min(\pi(x)q(y|x), \pi(y)q(x|y)) dx dy \\ &= [\pi(y)dy][q(x|y)a(y,x)dx] \\ &= \pi(dy)P(dx|y)\end{aligned}$$

If $x = y$, then the equation is trivial.

Metropolis-Hastings: Main properties

The Metropolis-Hastings (as described above) :

- ▶ is reversible and hence admits stationary distribution $\pi(d\cdot)$.
- ▶ is irreducible if

$$q(y|x) > 0, \text{ for every } x \in \mathcal{X}, y \in \mathcal{X}$$

since every set of \mathcal{X} can be reached in a single step

- ▶ is aperiodic, if it allows events $\{X^{(t+1)} = X^{(t)}\}$,
i.e. the probability of such an event is not zero

Metropolis-Hastings: Advantages/Challenges

Advantages:

We only need to know density $\pi(\cdot)$ up to a normalisation constant

$$a(x, x') = \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x|x')}{q(x'|x)}\right)$$

Challenges:

If the proposal distribution $q(d \cdot | \cdot)$ is poorly chosen:

- ▶ the exploration of the sampling space will be slow
- ▶ the standard error of the MC estimates will be large; because of high autocorrelations; $\tau_g = 1 + 2 \sum_{k=1}^{\infty} \text{Cor}(x_h, x_{h+k})$
- ▶ E.g. the number of rejections can be high

Metropolis-Hastings: Special cases

Popular special cases of the Metropolis-Hastings algorithm are:

IMH: The independence Metropolis-Hastings sampler

RWM: The Random Walk Metropolis algorithm*

MALA: The Langevin adjusted Hastings algorithm

... just different proposal distributions

Independence Metropolis-Hastings algorithm (IMH)

The proposal distribution $q(\cdot, d\cdot)$ is independent on the current state i.e. $q(x'|x) = q(x')$

$$X^{(n+1)} \sim P^{(\text{IMH})}(d\cdot | X^{(n)})$$

Given that the current state of the Markov chain is at $X^{(n)} = x$:

1. Generate x' from $q(d\cdot)$
2. Calculate

$$\begin{aligned} a(x, x') &= \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x|x')}{q(x'|x)}\right) \\ &= \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x)}{q(x')}\right) \end{aligned}$$

3. With probability $a(x, x')$ accept and set $X^{(n+1)} = x'$; otherwise reject and set $X^{(n+1)} = x$.

Independence Metropolis-Hastings: notes

1. The proposal distribution $q(d \cdot | \cdot)$ is independent on the current state; i.e. $q(x'|x) = q(x')$
2. The proposal distribution should be as close as possible to the target (stationary) distribution; i.e. $q(\cdot) \approx \pi(\cdot)$
3. Ideally, if $q(\cdot) = \pi(\cdot)$, then $a(x, x') = 1$, and the algorithm reduces to i.i.d. sampling from $\pi(d \cdot)$
4. It is a little bit difficult to find 'good' $q(\cdot)$ s.t. $q(\cdot) \approx \pi(\cdot)$, however possible if you try hard...

E.g.: If $\pi(\cdot)$ is uni-modal, $q(\cdot)$ can be a multivariate Normal distribution $N(\mu_\pi, \Sigma_\pi)$ where μ_π, Σ_π are properly chosen.

Random walk Metropolis algorithm (RWM)

The proposal distribution $q(d \cdot | \cdot)$ is s.t.

$$q(x'|x) = N(x'|x, \sigma^2 \mathbb{I}), \text{ for } \sigma^2 > 0$$

$$X^{(n+1)} \sim P^{(\text{RWM})}(d \cdot | X^{(n)})$$

Given that the current state of the Markov chain is at $X^{(n)} = x$:

1. Generate $x' \sim N(x, \sigma^2 \mathbb{I})$

2. Calculate

$$\begin{aligned} a(x, x') &= \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x|x')}{q(x'|x)}\right) = \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{N(x|x', \sigma^2 \mathbb{I})}{N(x'|x, \sigma^2 \mathbb{I})}\right) \\ &= \min\left(1, \frac{\pi(x')}{\pi(x)}\right) \end{aligned}$$

3. With probability $a(x, x')$ accept and set $X^{(n+1)} = x'$;
otherwise reject and set $X^{(n+1)} = x$.

Random walk Metropolis: notes 1

- ▶ Rational: “Local” exploration of the sampling space, around the neighbourhood of $X_n = x$.

$$q(x'|x) : \quad x' = x + \sigma z \quad ; \quad z \sim N(0, 1)$$

- ▶ Move towards modes of $\pi(\cdot)$ more often than moving away from them
 - ▶ “Uphill moves” are all accepted w.p. $a(x, x') = 1$
 - ▶ “Downhill moves” may be accepted w.p. $a(x, x') < 1$, or rejected w.p. $1 - a(x, x')$
- ▶ Advantages:
 - ▶ RWM is flexible: the choice of $q(\cdot | \cdot)$ is simple
 - ▶ RWM uses the previously simulated value x at stage $X^{(n)}$ to generate the proposed value x' for stage $X^{(n+1)}$.

Random walk Metropolis: notes 2

- ▶ RWM achieves optimal performance, if the proposal scale σ^2 leads to expected acceptance prob. $\bar{a}_{\text{opt}} \approx 0.234$

...if the components of $x := (x_1, \dots, x_d)$ are independent.

...however this rule leads to satisfactory performance in general cases

- ▶ Variations: $q(d \cdot | \cdot)$ can be any symmetric dist. s.t.

$$q(x'|x) = q(|x - x'|)$$

E.g. $q(d \cdot | \cdot)$:

propose $x' \sim U(x - \sigma, x + \sigma)$.

propose $x' = x + \sigma z$; $z \sim U(-1, 1)$

Langevin adjusted Hastings algorithm (MALA)

The proposal distribution $q(d \cdot | \cdot)$ is s.t.

$$q(x'|x) = N(x'|x + \frac{\sigma^2}{2} \nabla \log(\pi(x)), \sigma^2 \mathbb{I}), \text{ for } \sigma^2 > 0$$

$$X^{(n+1)} \sim P^{(\text{MALA})}(d \cdot | X^{(n)})$$

Given that the current state of the Markov chain is at $X^{(n)} = x$:

1. Generate $x' \sim N(x + \frac{\sigma^2}{2} \nabla \log(\pi(x)), \sigma^2 \mathbb{I})$
2. Calculate

$$\begin{aligned} a(x, x') &= \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{q(x|x')}{q(x'|x)}\right) \\ &= \min\left(1, \frac{\pi(x')}{\pi(x)} \frac{N(x|x' + \frac{\sigma^2}{2} \nabla \log(\pi(x')), \sigma^2 \mathbb{I})}{N(x'|x + \frac{\sigma^2}{2} \nabla \log(\pi(x)), \sigma^2 \mathbb{I})}\right) \end{aligned}$$

3. With probability $a(x, x')$ accept and set $X^{(n+1)} = x'$; otherwise reject and set $X^{(n+1)} = x$.

Langevin adjusted Hastings: notes

1. Goal: Direct the proposed values toward areas where density $\pi(\cdot)$ is likely to be larger by using information from $\pi(\cdot)$.
2. Rational: the inclusion of $\nabla \log(\pi(\cdot))$ in the proposal centre encourages moves towards the modes of $\pi(\cdot)$

$$q(x'|x) : x' = x + \frac{\sigma^2}{2} \nabla \log(\pi(x)) + \sigma z \quad ; \quad z \sim N(0, 1)$$

3. In difficult settings, exact gradients $\nabla \log(\pi(\cdot))$ can be replaced by numerical derivatives
4. MALA achieves optimal performance, if the proposal scale σ^2 leads to expected acceptance prob. $\bar{a}_{\text{opt}} \approx 0.57$

...if the components of $x := (x_1, \dots, x_d)$ are independent.

...however this rule leads to satisfactory performance in general cases

CODE

CODE

Tuning Metropolis-Hastings algorithms

The issue:

- ▶ How do we select a satisfactory proposal distr. $q(\cdot, d\cdot)$?
- ▶ Well, ... this is not easy in general Metropolis-Hastings algorithm
- ▶ But, ... for some special cases, it is possible by adjusting the proposals

Recall that :

- ▶ About RWM, the σ^2 is unknown
 - ▶ RWM can achieve satisfactory performance, if the proposal scale σ^2 leads to acc. prob. $a_{\text{opt}} \approx 0.234$
- ▶ About MALA, the σ^2 is unknown
 - ▶ MALA can achieve satisfactory performance, if the proposal scale σ^2 leads to acc. prob. $a_{\text{opt}} \approx 0.57$

An adaptive scheme for RWM, MALA

Goal: Adjust the proposal scaling σ^2 in RWM or MALA algorithms

For $n = 0, 1, 2, \dots$, iterate:

1. Simulate $X^{(n+1)}$ from $P_{\sigma_n^2}^{(\text{RWM})}(d \cdot |X^{(n)})$
2. Adjust σ^2 s.t. $\log(\sigma_{n+1}^2) = \log(\sigma_n^2) + \gamma_{n+1}(a_{n+1}^{\text{RWM}} - \bar{a}_{\text{opt}})$

Acceptance prob. of RWM at n -th iteration: a_n^{RWM}

Optimal acc. prob. value:
$$a_{\text{opt}} = \begin{cases} 0.234 & , \text{ for RWM} \\ 0.57 & , \text{ for MALA} \end{cases}$$

Gain sequence: $\gamma_n : \mathbb{N} \rightarrow \mathbb{R}_+$, a decreasing function $\gamma_n \searrow 0$

E.g. $\gamma_n = C/n^\varsigma$, $C > 0$, $\varsigma \in (0.5, 1)$

Adaptive RWM and MALA algorithms: notes 1

Gain sequence γ_n

- ▶ γ_n must present a smooth slow decay
- ▶ As $n \uparrow$, $\gamma_n \downarrow$, and the influence of adaptation vanishes
- ▶ A reasonable choice is

$$\gamma_n = C/n^\varsigma, \quad C > 0, \quad \varsigma \in (0.5, 1)$$

- ▶ ς controls the speed that γ_n decays to 0

Christophe Andrieu and Johannes Thoms (2008). A tutorial on adaptive MCMC

An adaptive scheme for RWM, MALA: notes 2

Rational: At state n ,

- ▶ if $a_{n+1}^{\text{RWM}} < \bar{a}_{\text{opt}}$,
 - $\implies \log(\sigma_{n+1}^2) < \log(\sigma_n^2)$
 - $\implies \sigma_n^2$ decreases
 - \implies RWM/MALA will perform smaller steps at stage $n + 1$
- ▶ if $a_{n+1}^{\text{RWM}} > \bar{a}_{\text{opt}}$,
 - $\implies \log(\sigma_{n+1}^2) > \log(\sigma_n^2)$
 - $\implies \sigma_n^2$ increases
 - \implies RWM/MALA will perform larger steps at stage $n + 1$

Christophe Andrieu and Johannes Thoms (2008). A tutorial on adaptive MCMC

CODE

CODE

Blockwise MCMC samplers

Consider r.v. $X := (X_1, \dots, X_d)$ that follows $X \sim \pi(dX)$.

Challenge: In many cases, it is difficult to select appropriate proposals to construct an efficient Metropolis-Hastings algorithm that ‘updates’ simultaneously the whole $X := (X_1, \dots, X_d)$

Reasons: X can be high dimensional.

Different X_i may have different ranges, types, etc.

etc...

Cure: ‘Break’ sampling of X by combining M-H algorithms targeting the conditional distributions of $\pi(d\cdot)$

Blockwise MCMC sampler

Consider r.v. $X := (X_1, \dots, X_d)$ that follows $X \sim \pi(dX)$.

How to generate $X^{(n)} \sim P^{(\text{blockwise})}(d \cdot | X^{(n-1)})$ targeting $\pi(dX)$??

At iteration n :

- ▶ Simulate $X_1^{(n)} \sim P_1^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_1^{(n)} | X_2^{(n-1)}, \dots, X_d^{(n-1)})$
 \vdots
- ▶ Simulate $X_i^{(n)} \sim P_i^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_i^{(n)} | X_1^{(n)}, \dots, X_{i-1}^{(n)}, X_{i+1}^{(n-1)}, \dots, X_d^{(n-1)})$
 \vdots
- ▶ Simulate $X_d^{(n)} \sim P_d^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_d^{(n)} | X_1^{(n)}, X_3^{(n)}, \dots, X_{d-1}^{(n)})$

Set $X^{(n)} = (X_1^{(n)}, X_3^{(n)}, \dots, X_d^{(n)})$

For example, for the i -th block

Simulate $X_i^{(n)} \sim P_i^{(\text{RWM})}(\mathbf{d} \cdot | \cdot)$ targeting
 $\pi(\mathbf{d}X_i^{(n)} | X_1^{(n)}, \dots, X_{i-1}^{(n)}, X_{i+1}^{(n-1)}, \dots, X_d^{(n-1)})$ (via RWM)

1. Generate $x' \sim \mathcal{N}(x_i^{(n-1)}, \sigma^2 \mathbb{I})$

2. Calculate

$$\begin{aligned} a(x, x') &= \min\left(1, \frac{\pi(x' | x_1^{(n)}, \dots, x_{i-1}^{(n)}, x_{i+1}^{(n-1)}, \dots, x_d^{(n-1)})}{\pi(x_i^{(n-1)} | x_1^{(n)}, \dots, x_{i-1}^{(n)}, x_{i+1}^{(n-1)}, \dots, x_d^{(n-1)})}\right) \\ &= \min\left(1, \frac{\pi(x_1^{(n)}, \dots, x_{i-1}^{(n)}, x', x_{i+1}^{(n-1)}, \dots, x_d^{(n-1)})}{\pi(x_1^{(n)}, \dots, x_{i-1}^{(n)}, x_i^{(n-1)}, x_{i+1}^{(n-1)}, \dots, x_d^{(n-1)})}\right) \end{aligned}$$

3. With probability $a(x_i^{(n-1)}, x')$ accept and set $X_i^{(n)} = x'$;
otherwise reject and set $X_i^{(n)} = x_i^{(n-1)}$.

Gibbs sampler (special case of Blockwise MCMC sampler)

Consider r.v. $X := (X_1, \dots, X_d)$ that follows $X \sim \pi(dX)$.

If all the full conditional dist of $\pi(dX)$ can be sampled directly

How to $X^{(n)} \sim P^{(\text{Gibbs})}(d \cdot | X^{(n-1)})$ targeting $\pi(dX)$??

At iteration n :

- ▶ Simulate $X_1^{(n)} \sim \pi(dX_1^{(n)} | X_2^{(n-1)}, \dots, X_d^{(n-1)})$
⋮
- ▶ Simulate $X_i^{(n)} \sim \pi(dX_i^{(n)} | X_1^{(n)}, \dots, X_{i-1}^{(n)}, X_{i+1}^{(n-1)}, \dots, X_d^{(n-1)})$
⋮
- ▶ Simulate $X_d^{(n)} \sim \pi(dX_d^{(n)} | X_1^{(n)}, X_3^{(n)}, \dots, X_{d-1}^{(n)})$

Set $X^{(n)} = (X_1^{(n)}, X_3^{(n)}, \dots, X_d^{(n)})$

CODE

CODE

Blockwise MCMC sampler: notes

The blockwise MCMC sampler admits $\pi(dX)$ as stationary distribution

Systematic sweep: (described above)

- ▶ The blocks are updated in a fix order:
$$P^{(\text{blockwise})} = P_1^{(M-H)} P_2^{(M-H)} \dots P_d^{(M-H)}$$
- ▶ The Markov chain is NOT reversible

Permutation sweep:

- ▶ The blocks are updated in a random permutation order p
$$P^{(\text{blockwise})} = P_{p(1)}^{(M-H)} P_{p(2)}^{(M-H)} \dots P_{p(d)}^{(M-H)}$$
- ▶ The Markov chain is reversible

Random sweep:

- ▶ At each iteration, randomly select and update ONLY one block
$$P^{(\text{blockwise})} = \frac{1}{d} \sum_{i=1}^d P_i^{(M-H)}$$
- ▶ The Markov chain is reversible

Blockwise MCMC sampler (Permutation sweep)

Consider r.v. $X := (X_1, \dots, X_d)$ that follows $X \sim \pi(dX)$.

How to generate $X^{(n)} \sim P^{(\text{blockwise})}(d \cdot | X^{(n-1)})$ targeting $\pi(dX)$??

At iteration n :

- ▶ Generate a random permutation $p = (p(1), \dots, p(d))$
 - ▶ Simulate $X_{p(1)}^{(n)} \sim P_{p(1)}^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_{p(1)}^{(n)} | \text{all the rest})$
 - ▶ Simulate $X_{p(i)}^{(n)} \sim P_{p(i)}^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_{p(i)}^{(n)} | \text{all the rest})$
 - ▶ Simulate $X_{p(d)}^{(n)} \sim P_{p(d)}^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_{p(d)}^{(n)} | \text{all the rest})$

Set $X^{(n)} = (X_1^{(n)}, X_3^{(n)}, \dots, X_d^{(n)})$

Blockwise MCMC sampler (Random sweep)

Consider r.v. $X := (X_1, \dots, X_d)$ that follows $X \sim \pi(dX)$.

Generate $X^{(n)} \sim P^{(\text{blockwise})}(d \cdot | X^{(n-1)})$ targeting $\pi(dX)$

At iteration n :

- ▶ Select block $i \sim U\{1, \dots, d\}$, at random
- ▶ Simulate $X_i^{(n)} \sim P_i^{(\text{M-H})}(d \cdot | \cdot)$ targeting $\pi(dX_i^{(n)} | \text{all the rest})$

Set $X^{(n)} = (X_1^{(n-1)}, \dots, X_{i-1}^{(n-1)}, X_i^{(n)}, X_{i+1}^{(n-1)}, \dots, X_d^{(n-1)})$

Improving the quality of the MCMC sample

After we generate the MCMC sample $\{X_1, X_2, X_3, \dots\}$

Burn-in: Use only the generated Markov chain in the stationarity

- ▶ Discard the first few iterations (e.g. first b steps) of the Markov chain as a burn-in period
- ▶ Keep only the last tail of the Markov chain

E.g. keep $\tilde{X} = \{X_b, X_{b+1}, X_{b+2}, X_{b+3}, \dots\}$

Thinning: Try to reduce the autocorrelation by sub-sampling

- ▶ Use only every k -th element of the generated Markov chain

E.g. use $\tilde{\tilde{X}} = \{\tilde{X}_1, \tilde{X}_{1+k}, \tilde{X}_{1+2k}, \dots\} = \{X_{b+1}, X_{b+k}, X_{b+2k}, \dots\}$
... k -step thinning

Application: Inference on what?

Compute densities:

- ▶ $\pi(\alpha, \beta | D)$
- ▶ $\pi(\alpha | D) = \int \pi(\alpha, \beta | D) d\beta,$
- ▶ $\pi(\beta | D) = \int \pi(\alpha, \beta | D) d\alpha$

Compute expectations:

- ▶ $E(\alpha | D) = \int \alpha \pi(\alpha | D) d\alpha$
- ▶ $E(\beta | D) = \int \beta \pi(\beta | D) d\beta$
- ▶ $\Pr(t = 66.0 | D) = E(p(t = 66.0 | (\alpha, \beta)) | D)$
$$= \int p(t = 66.0 | (\alpha, \beta)) \pi(\alpha, \beta | D) d\alpha d\beta$$

Application: How to facilitate inference?

Generate sample: $\{(\alpha_n, \beta_n) \sim P(d \cdot | (\alpha_{n-1}, \beta_{n-1})); n = 1, \dots, N\}$

Estimate densities:

- ▶ $\hat{\pi}(\alpha, \beta | D)$: with the histogram of $\{(\alpha_n, \beta_n); n = 1, 2, \dots, N\}$
- ▶ $\hat{\pi}(\alpha | D)$: with the histogram of $\{\alpha_n; n = 1, 2, \dots, N\}$
- ▶ $\hat{\pi}(\beta | D)$: with the histogram of $\{\beta_n; n = 1, 2, \dots, N\}$

Estimate expectations:

- ▶ $\widehat{E(\alpha | D)} = \bar{\alpha}^{(N)} = \frac{1}{N} \sum_{n=1}^N \alpha_n$, with $\text{s.e.}(\widehat{E(\alpha | D)}) = \sqrt{\frac{1}{N} s_{\alpha}^2 \tau_{\alpha}}$
- ▶ $\widehat{E(\beta | D)} = \bar{\beta}^{(N)} = \frac{1}{N} \sum_{n=1}^N \beta_n$, with $\text{s.e.}(\widehat{E(\beta | D)}) = \sqrt{\frac{1}{N} s_{\beta}^2 \tau_{\beta}}$
- ▶ $\widehat{\text{Pr}}(t = 66.0 | D) = \bar{p}^{(N)} = \frac{1}{N} \sum_{n=1}^N \overbrace{p(t = 66.0 | (\alpha_n, \beta_n))}^{=p_n}$

with $\text{s.e.}(\widehat{\text{Pr}}(t = 66.0 | D)) = \sqrt{\frac{1}{N} s_p^2 \tau_p}$